

Reservoir Sampling

Anže Pečar, Miha Zidar

Abstract—The abstract goes here.[TODO]

Index Terms—[TODO]

I. INTRODUCTION

THE problem with random sampling is to select a random sample of size n from a set of size N . Many algorithms have been developed for this problem when the value of N is known beforehand. Some problems, that we encounter in the real world, do not have a specified N or N cannot be determined efficiently. Those kind of problems will be the focus of this paper. We shall take a look at older algorithms such as Algorithm R which was developed to efficiently and accurately create a random sample from a tape in one pass, as well as newer algorithms, such as [TODO]

Many reservoir algorithms make the assumption that the data does not change over time. Algorithm R, for an example, samples starting data with a greater frequency than the data at the end of the tape, which is efficient but it fails to react to emerging trends in the data. If we are interested in trends in our data a different kind of algorithm needs to be used. [TODO]

December 30, 2011

II. ALGORITHM R AND ITS IMPROVEMENTS

A. Motivation

In order to better understand the complex algorithms in use today we first need to understand the idea behind simpler algorithms used in the past. In this section we shall describe Algorithm R and its improvement - Algorithm Z.

B. Algorithm R

Algorithm R is a reservoir algorithm written by Alan Waterman. It works as follows: we fill up our reservoir of n records with the first n records of the file we are processing. The $t + 1$ st record then has a $n/(t + 1)$ chance of being in our reservoir of size n . The candidate it replaces is chosen randomly from the n candidates.

Below is the implementation of Algorithm R in *Python*:

```
from random import randrange

N = [0] * n # initial reservoir
for i in xrange(n):
    N[i] = READ_NEXT_RECORD()

t = n
while EXISTS_NEXT_RECORD():
```

```
t += 1
M = randrange(0, t, 1)
if M < n:
    N[M] = READ_NEXT_RECORD()
else:
    READ_NEXT_RECORD()
```

Explanation:

The list N is our reservoir in which we store a random sample of records. We start off with an empty list of length n in which we store the first n elements (the for loop). The function `READ_NEXT_RECORD()` returns the next record in the stream. After we have inserted the first n elements into our reservoir, we enter the *while* clause, which runs as long as there are still records in the stream. In each run of the *while* clause we first increment the counter of records t and then calculate a random number between 0 and t . We store the random number in M . If M is lower than n we store the record at index M , otherwise we skip the record. [AM I EXPLAINING THE OBVIOUS?]

C. Algorithm Z

It turns out we do not need to go over every single record in order to get a random set in the reservoir. We can skip a random number of records each time just as long as the probability of a record being in the reservoir does not change. This is the idea behind Algorithm Z.

Below is the first part of Algorithm Z implemented in *Python*:

```
from random import randrange

N = [0] * n # initial reservoir
for i in xrange(n):
    N[i] = READ_NEXT_RECORD()

t = n
thresh = T * n
num = 0
while EXISTS_NEXT_RECORD() and t <= thresh:
    t += 1
    num += 1
    V = random()
    S = 0
    quot = num/t
    while quot > V:
        S += 1
        t += 1
        num += 1
        quot = (quot * num)/t

    SKIP_RECORDS(S)

if EXISTS_NEXT_RECORD():
    M = randrange(0, n)
    N[M] = READ_NEXT_RECORD()
```

The algorithm works best if T is between 10 and 40. When the *thresh* value gets large enough the algorithm starts using the rejection technique. The rejection technique generates S even more slowly than in the above algorithm, which makes the algorithm faster as there are even more records that we are able to skip. The main difference between *Algorithm R* and *Algorithm Z* is that the former computes every single record while the latter skips as many records as possible.



Anže Pečar Biography text here.

III. ONLINE LEARNING

[TODO]

IV. CONCLUSION

The conclusion goes here.

APPENDIX A APPENDIX TITLE

Appendix one text goes here.

APPENDIX B

Appendix two text goes here.

ACKNOWLEDGMENT

The authors would like to thank...

REFERENCES

- [1] H. Kopka and P. W. Daly, *A Guide to L^AT_EX*, 3rd ed. Harlow, England: Addison-Wesley, 1999.



Miha Zidar Biography text here.